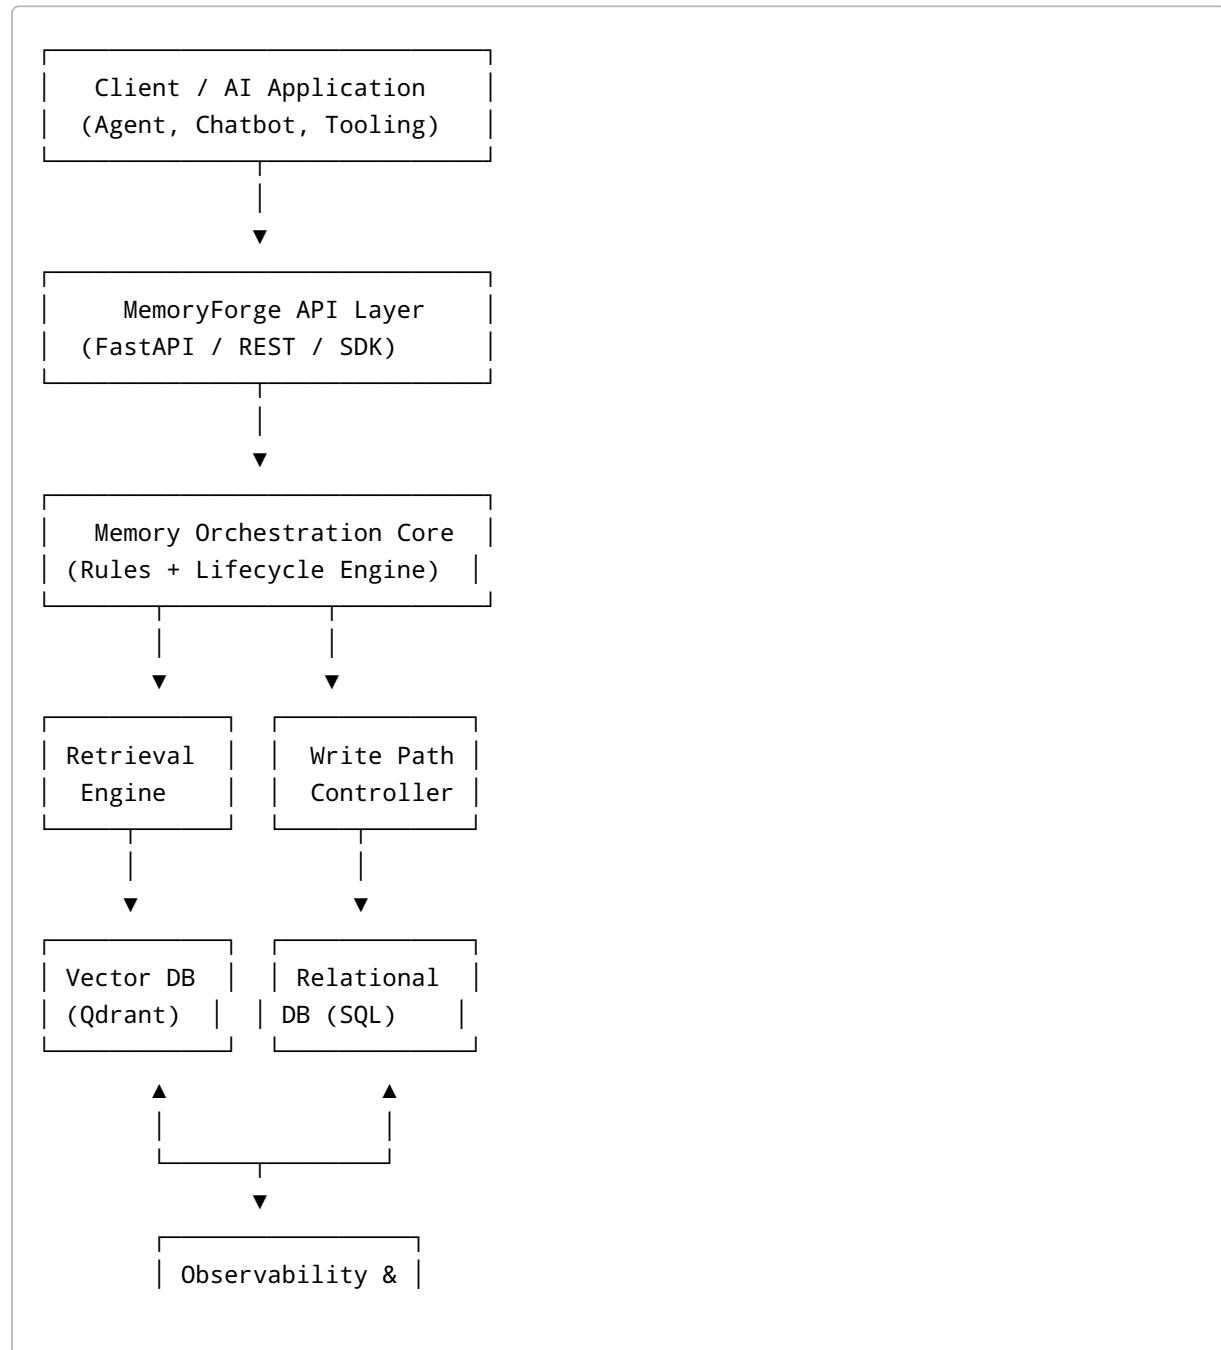


Smriti - System Architecture (Production-Ready)

This architecture is designed to be **buildable, debuggable, and trustworthy**. It avoids premature scale while enforcing correctness, memory control, and explainability.

1. High-Level Architecture (Mental Model)



2. Core Architectural Principles (Non-Negotiable)

1. Rules before LLMs
2. Memory is versioned, never overwritten
3. Reads are cheap, writes are expensive
4. Every decision is explainable
5. Failure is visible, not silent

3. Component Breakdown

3.1 API Layer

Responsibilities:

- Authentication / namespace isolation
- Input validation
- Rate limiting
- SDK abstraction

Tech:

- FastAPI
- Pydantic schemas

No business logic here.

3.2 Memory Orchestration Core (Heart of V1)

This is the system's brain.

Responsibilities:

- Enforce memory lifecycle rules
- Decide read vs write vs update
- Route to correct memory type
- Enforce confidence thresholds

Important: This layer does NOT store data. It controls flow.

3.3 Memory Write Path (Controlled & Defensive)

```
Input → Validation → Classification → Rule Check → LLM Assist → Persist → Audit
```

Key rules:

- LLM suggestions are advisory
- Conflicts create new versions
- Deletes require explicit approval

Writes are slow by design.

3.4 Retrieval Engine (Context-Aware)

Capabilities:

- Memory-type filtering
- Temporal constraints
- Active/inactive filtering
- Relevance thresholding

Output:

- Ranked, minimal memory set
- Token-budget-aware

This replaces naive RAG.

4. Memory Types & Storage Mapping

Memory Type	Storage	Mutability
Working	In-memory + SQL	Mutable
Episodic	SQL	Immutable
Semantic	SQL + Vector	Versioned
Metadata	SQL	Mutable

Graph DB is intentionally excluded in V1.

5. Storage Layer

5.1 Relational DB (Source of Truth)

Stores:

- Memory records
- Versions
- Confidence scores
- Status (active/inactive)
- Source & timestamps

Why SQL:

- Strong consistency
 - Easy audits
 - Human debuggable
-

5.2 Vector DB (Similarity Only)

Stores:

- Embeddings for semantic memory

Used ONLY for:

- Candidate recall

Never used for authority.

6. LLM Interaction Layer

LLMs are used for:

- Summarization
- Classification
- Suggesting memory updates

LLMs are NOT used for:

- Deletion
- Final conflict resolution
- Truth authority

All LLM outputs must include confidence.

7. Contradiction Handling Flow

```
New Fact → Similarity Search → Conflict Detection  
→ Version Creation  
→ Resolution Policy Applied  
→ One Active, Others Inactive
```

Nothing is lost.

8. Observability & Trust Layer

Every memory event logs:

- Input
- Decision path
- Rules triggered
- LLM output
- Final action

This is critical for production trust.

9. Deployment Model (V1)

Supported:

- Single-node deployment
- Docker-based
- Local-first

Explicitly NOT required:

- Kubernetes
 - Kafka
 - GPUs
-

10. Why This Architecture Will Work

Because it:

- Minimizes moving parts

- Prioritizes correctness over speed
 - Treats memory as state, not text
 - Gives humans control
-

Senior Engineer's Rule of Thumb

If you can't explain why a memory exists, the system is already broken.

This architecture prevents that.